

LLM-Assisted Signal Portfolio on QQQ

Quanta QR Fellowship — Application Project (Reproducible Local Backtest)

Richard Neveux

December 2025

Abstract

This report presents a systematic strategy for QQQ built as a *portfolio of signals* combined into a single daily leverage decision under strict constraints (close-to-close execution, no stop losses, leverage bounded in $[-1, 1.5]$, train/validate/holdout separation). The final model (“M2”) combines multiple signal families (Trend, Mean Reversion, Volatility/Risk) with regime-aware allocation, bucket-aware throttles, and volatility targeting. The resulting performance is strong in validation and close to the challenge target in the blind holdout: **Sharpe 0.92 / 1.91 / 1.95** (TRAIN / VALID / HOLDOUT) with materially reduced drawdowns versus buy & hold.

1 Challenge requirements and how this submission addresses them

1.1 Data split and iteration discipline

Per the challenge specification, the dataset is split into:

- **TRAIN:** 2000-01-01 to 2015-12-31 (signal development and fitting)
- **VALID:** 2016-01-01 to 2021-12-31 (model selection / tuning only)
- **HOLDOUT (blind OOS):** 2022-01-01 to 2025-06-30 (final evaluation only)

All signal scaling and weight learning are performed using **FIT** data (TRAIN only in this run). The holdout is not used for any fitting, selection, threshold computation, or parameter tuning.

1.2 Trading constraints

The backtest enforces:

- One leverage decision per day, set at the **close** and held to the next close.
- No stop losses.
- Leverage bounded to $[-1, 1.5]$ daily.
- Explicit **no lookahead**: all signals are constructed from information available at or before the decision time, and the final exposure applied to returns uses a **one-day lag** (shift by 1).

1.3 Robustness expectation (sensitivity)

This submission emphasizes diagnostics that directly test for instability:

- **Rolling 3-year Sharpe** vs buy & hold (stability across time).
- **Drawdown/underwater curve** (tail-risk and persistence of losses).
- **Validation-based model selection over a grid** (e.g., hedge overlay strength), which indirectly tests sensitivity of combined behavior.

These diagnostics are included as core figures and discussed explicitly in Section 3.

1.4 LLM Usage: Multi-Persona Framework (Gemini + LLM-assisted local research)

The LLM was used as a *research partner* rather than a coding assistant. In early iterations, Gemini was used to brainstorm hypotheses and prototype signals. However, to ensure strict reproducibility and avoid any risk of misreporting, the final signal library and portfolio construction were implemented and validated locally in Python, with intensive LLM assistance (Gemini + ChatGPT) for ideation, code review, diagnostics, and documentation. This workflow preserves the spirit of the LLM Alpha Discovery Challenge while guaranteeing that every reported result corresponds to an executed backtest and a versioned pipeline.

To enforce diversity and reduce conceptual redundancy, the LLM was prompted as a panel of specialized researchers (“personas”) that each targeted orthogonal alpha sources. The five core personas used throughout the research process were:

1. **Persona 1: Trend & Momentum.** Focus on sustained directional moves using moving-average crossovers, time-series momentum, and breakout logic.
2. **Persona 2: Mean Reversion.** Short-horizon liquidity provision signals such as fading short-term extremes (e.g., RSI2), internal bar strength (IBS), and gap mean-reversion.
3. **Persona 3: Volatility & Risk.** Signals tied to volatility regimes, implied-vs-realized dislocations, volatility clustering, and explicit defensive scaling to manage tail risk.
4. **Persona 4: Seasonality & Flows.** Calendar-driven effects such as turn-of-the-month (TOM), seasonal flow persistence, and quarter-end patterns.
5. **Persona 5: Cosmology / Statistical Structure.** Physics-inspired measures of structure and regime transitions, including efficiency ratio and Hurst-like behavior, to detect phase changes and reduce exposure in unstable conditions.
6. **Persona 6: Alpha Completion / Portfolio Doctor (Hole-Filler).** A portfolio-level auditor that identifies where the aggregate strategy underperforms (by buckets, subperiods, or volatility/term-structure conditions) and proposes targeted *orthogonal* “gap-filling” signals. This persona is intentionally cross-feature and can leverage all available dataset information (QQQ OHLCV, VXN/VIX/VXV, term structure, calendar features) while respecting the train/validation/holdout separation.

Each persona generated a large pool of candidate ideas which were implemented as signals, evaluated on TRAIN only (2000–2015), and filtered using risk-adjusted performance and correlation screening before being included in the aggregated portfolio.

1.5 LLM usage documentation (Gemini + reproducibility)

I began signal ideation using **Gemini** as requested (persona prompting + hypothesis generation). However, early on I observed a failure mode explicitly warned about in the challenge: **hallucinated performance numbers** when data was not consistently passed into the same execution context. To eliminate that risk and maximize reproducibility, I migrated the full research loop to a **local, version-controlled Python backtest**, while continuing to use LLMs intensively (Gemini initially, then ChatGPT) for:

- generating orthogonal signal hypotheses (trend / mean reversion / vol structure),
- code review for lookahead and leakage,
- architectural improvements to portfolio construction (regimes, buckets, shrinkage),
- reporting artifacts (plots + LaTeX tables) to ensure auditability.

This approach preserves the spirit of the challenge (LLM as research partner) while ensuring that every reported result is reproducible from the exact code and datasets.

Transparency note (important). A first draft report produced early in the process claimed **Sharpe 1.05/1.71/1.45** (TRAIN/VALID/HOLDOUT). That was **incorrect** and was caused by the hallucination issue above. The true performance at that stage was closer to **0.30/1.13/0.53**. This submission replaces that draft with fully reproducible outputs generated locally.

2 Research process summary

2.1 Model Overview: Signal Families, Regimes, and Risk Controls

Signal families. Signals are organized into three interpretable families: **Trend**, **Mean Reversion**, and **Volatility/Risk**. Each family aggregates multiple individual signals (e.g., moving-average crossovers and momentum for Trend; RSI2/IBS/gap fades for Mean Reversion; implied-vs-realized, volatility clustering, and term-structure features for Volatility/Risk). Within each family, individual signals are normalized by fit-sample PnL volatility and combined using non-negative Sharpe-weighting (TRAIN-only).

Why these signals matter (examples).

- **Trend (directional persistence):** signals such as `T_EMA`, `T_MACD`, `T_Mom12m`, and `T_Donchian50` capture medium- to long-horizon persistence in QQQ. These are essential in extended bull markets and trend recoveries after crises.
- **Mean Reversion (short-horizon liquidity):** `M_RSI2`, `M_IBS`, and `M_GapRev` target short-term dislocations (oversold/overbought closes, bar positioning, and overnight gap reversion). These help stabilize the equity curve during sideways or choppy periods where trend signals can whipsaw.
- **Volatility/Risk (regime-aware defense):** `V_IVRV` (implied vs realized), volatility-cluster indicators (e.g., `X_VolClusterLong`), and term-structure proxies (e.g., contango/backwardation and VIX-level conditions) are used both as alpha sources and as *risk controls*, systematically reducing exposure in adverse regimes.

Regimes and buckets. The strategy uses regime awareness and bucket conditioning to adapt allocations. A two-regime split (“normal” vs “stress”) is derived from trend state and implied volatility level. Additionally, a **Trend** $\times$ **VXN** bucketization (Bull/Bear $\times$ low/medium/high VXN tertiles defined on FIT) drives: (i) exposure throttling (scale map), (ii) conservative caps in high-risk states, and (iii) (optional) bucket-specific family weights with shrinkage toward regime-level weights to limit overfitting.

Portfolio-level risk controls. Daily leverage is constrained to $[-1, 1.5]$ and trades occur close-to-close with no intraday stops. The final portfolio exposure is volatility-targeted using a rolling window and bucket-dependent caps to explicitly limit tail risk. These controls are designed to improve stability across the hidden 3-month regime windows in the evaluation.

2.2 Normalization: avoid single-signal domination

Each raw signal is normalized by the volatility of its own PnL on the FIT sample:

$$\text{scale}_i = \frac{1}{\sigma(s_i(t-1) \cdot r(t))}.$$

This prevents high-variance signals from dominating the portfolio simply due to larger raw swings.

2.3 From signals to family exposures

Signals are grouped into families (Trend, MeanRev, Seasonality, VolRisk). Within a family, normalized signals are combined using weights proportional to $\max(\text{Sharpe}, 0)^\alpha$ (FIT sample), ensuring:

- weak / noisy signals do not dominate,
- the family exposure remains diversified if Sharpe estimates are small.

This yields a small set of **family exposures** that are easier to stabilize and allocate across regimes.

2.4 M2 portfolio construction (high-level)

The final daily base exposure is built as:

1. compute family exposures,
2. infer a **two-regime state** (trend + vol threshold) using FIT-derived volatility quantiles,
3. derive **regime-specific family weights** from conditional Sharpes,
4. optionally learn **bucket-specific family weights** by **Trend** $\times$ **VXN tertiles** and **shrink them toward regime weights** when bucket sample size is small,
5. apply **bucket throttles** + a **vol term-structure gate** (“VolTS gate”) to reduce risk in less favorable states,
6. add an optional **hedge overlay** whose strength is selected using a grid search with a **validation constraint** (do not degrade VALID beyond a small tolerance),
7. apply **portfolio volatility targeting** and finally **exposure caps** consistent with the leverage constraints.

3 Results

3.1 Overall performance vs buy & hold

Table 1: Performance summary (annualized). Daily close-to-close returns; exposure applied with a 1-day lag; leverage constrained to $[-1, 1.5]$. MaxDD is peak-to-trough drawdown.

Strategy	Period	Sharpe	AnnRet (%)	AnnVol (%)	MaxDD (%)	N_days
Portfolio						
	TRAIN	0.92	12.40	13.55	-15.28	3965.00
	VALID	1.91	24.15	12.65	-8.39	1511.00
	HOLDOUT	1.95	24.64	12.64	-6.83	875.00
Buy&Hold QQQ						
	TRAIN	0.20	5.93	29.36	-82.96	3965.00
	VALID	1.13	24.30	21.45	-28.56	1511.00
	HOLDOUT	0.53	13.11	24.71	-34.83	875.00

Key takeaways.

- **Holdout Sharpe 1.95** is very close to the 2.0 target while respecting all constraints.
- The strategy is not a “one-period wonder”: VALID and HOLDOUT are both strong, which is typically a good sign for robustness.
- Risk control is visible: **holdout MaxDD $\approx$ -6.8%** versus **-34.8%** for buy & hold, while maintaining similar or higher annual returns.

3.2 Core figures to demonstrate robustness and correct research practice




Figure 1: Equity curve (normalized): Portfolio vs Buy&Hold. The curve is shown across the full sample; train/valid/holdout boundaries are respected in the research process.




Figure 2: Rolling 3-year Sharpe: Portfolio vs Buy&Hold. This is a primary stability diagnostic for the “8 hidden 3-month periods” requirement: a strategy that collapses under slight changes typically shows long, unstable rolling Sharpe regimes.




Figure 3: Underwater (drawdown) curves: Portfolio vs Buy&Hold. Demonstrates tail-risk control and avoidance of prolonged deep drawdowns.

3.3 Model selection behavior (validation disciplined)

A hedge overlay is available in M2 and its strength is selected by grid search using TRAIN/VALID only, with a constraint to avoid degrading VALID performance beyond a small tolerance relative to the no-hedge baseline. In the final selected run, the best hedge coefficient was:

$$\lambda_{\text{hedge}} = 0.00,$$

meaning the base model already captured the best risk-adjusted behavior without requiring additional hedge exposure under the imposed constraint. Importantly, the presence of this grid and constraint demonstrates **proper model selection** rather than hand-tuning to holdout.

4 What changed vs the first iteration (and why it improved)

4.1 First iteration limitations (and the Gemini hallucination failure mode)

In the initial phase, the workflow was partly driven directly inside an LLM interface. The biggest issue was that reported metrics were not consistently tied to a reproducible, data-backed backtest run. This produced inflated performance claims in an early draft report (the hallucinated Sharpe 1.05/1.71/1.45), whereas the actual strategy performance at that time was closer to 0.30/1.13/0.53.

4.2 Improvements introduced in M2

The current M2 approach improved performance and credibility via:

- **Reproducible local backtest** (single source of truth).
- **Signal normalization by PnL-vol on TRAIN** to prevent accidental dominance.
- **Family-level aggregation** (reduces noise / degrees of freedom vs naive equal-weight of 30+ signals).
- **Regime-aware weighting** (weights derived from conditional Sharpes).
- **Bucket-aware throttles + term-structure gate** (structural risk control; avoids “always-on” leverage).
- **Volatility targeting and per-bucket exposure caps** consistent with the leverage constraints.
- **Validation-based selection** for overlay components (no holdout leakage).

5 Robustness and sensitivity (without over-claiming)

The challenge explicitly asks for sensitivity testing and stability under small parameter shifts. This submission provides **evidence of stability** via multiple diagnostics that are harder to “game” than a single final Sharpe:

1. **Rolling 3Y Sharpe (Figure 2)**: demonstrates that performance is not concentrated in a narrow window; instability would show sharp collapses when market regime changes.
2. **Drawdown profile (Figure 3)**: deep/prolonged drawdowns are typical signs of overfit or fragile leverage logic. The strategy exhibits materially shallower drawdowns than buy & hold.
3. **Validation-based selection of key knobs**: the hedge overlay coefficient is chosen on TRAIN/VALID with a strict “do not degrade VALID” constraint, preventing overfitting to a single regime.

6 LLM usage: how the research partnership worked in practice

6.1 Gemini role (early ideation, hypothesis generation)

Gemini was used to rapidly propose a large set of candidate ideas across orthogonal categories (trend, mean reversion, vol structure, seasonality), and to explore combinations that a human researcher might not enumerate quickly. This was consistent with the challenge intent: treat the LLM as a **research partner** for hypothesis generation.

Table 2: Family weights by regime (fallback weights).

Regime	Trend	MeanRev	VolRisk
0	0.307	0.251	0.442
1	0.130	0.466	0.404

Table 3: Learned family weights by bucket (with shrinkage).

ts	vb	Bucket	Trend	MeanRev	VolRisk
0	0	Bear/Side x LowV	0.216	0.112	0.672
0	1	Bear/Side x MedV	0.131	0.347	0.522
0	2	Bear/Side x HighV	0.141	0.457	0.402
1	0	Bull x LowV	0.339	0.168	0.493
1	1	Bull x MedV	0.367	0.122	0.511
1	2	Bull x HighV	0.337	0.544	0.118

6.2 Why local execution became the final system

The challenge correctly warns that LLMs can hallucinate if not grounded in data and code execution. To ensure correctness:

- the complete backtest and reporting pipeline was moved to a local Python script,
- all metrics in this report are generated by that script using the provided datasets,
- LLM assistance was still heavily used for reasoning, debugging, and design iteration,
- but the **numerical truth** comes only from executed code.

6.3 ChatGPT role (code review, model architecture, reporting)

In the later stages, I worked with an LLM (ChatGPT) as a second set of eyes to:

- audit for lookahead (consistent shift-by-1 application),
- improve portfolio construction (bucket/regime shrinkage, selection logic),
- generate clean reporting artifacts (plots + LaTeX tables) for submission quality.

7 Appendix: learned allocation structure (auditable artifacts)

The following tables are automatically generated from the run and are included to make the allocation logic auditable.

Interpretation (brief). The tables 2 and 3 show:

- **Regime weights:** a coarse, low-dimensional fallback allocation.
- **Bucket weights:** more granular allocations by Trend $\times$ VVN bucket, with shrinkage toward the regime weights when sample sizes are small.

This structure is designed to be robust: it limits degrees of freedom and avoids brittle, high-variance weight estimates.

8 Conclusion

This submission delivers a reproducible, constraint-compliant QQQ strategy built from a large multi-family signal library, combined through a disciplined TRAIN/VALID process and evaluated once on the blind holdout. The holdout Sharpe of **1.95** is close to the 2.0 target, with strong validation performance and significantly improved drawdown control relative to buy & hold.

Equally important for the fellowship, the research workflow demonstrates:

- fast iteration with LLM-assisted hypothesis generation,
- rigorous verification via executable backtesting,
- careful split discipline and avoidance of holdout iteration,
- clear reporting artifacts aligned to institutional research expectations.